

Searching help pages of R packages

by *Spencer Graves, Sundar Dorai-Raj, and Romain François*

The `sos` package provides a means to quickly and flexibly search the help pages of contributed packages, finding functions and datasets in seconds or minutes that could not be found in hours or days by any other means we know.

The main capability of this package is the `findFn` function, which scans the "function" entries in Jonathan Baron's "R site search" database and returns the matches in a `data.frame` of class `findFn` (Baron, 2009). Baron's is one of five search capabilities currently identified under "search" from the main "www.r-project.org" web site. It includes options to search the help pages of R packages contributed to CRAN (the Comprehensive R Archive Network) plus a few other publicly available packages as well as selected mailing list archives, primarily "R-help". The `findFn` function focuses only on the help pages in this database.

The print method for objects of class `findFn` displays the results as a table in a web browser with links to the individual help pages, sorted by package displaying the one with the most matches first. This is different from the `RSiteSearch` function, as `findFn` returns the results in R as a `data.frame`, which can be further manipulated, and the ultimate display in a web browser is a table, unlike the list produced by `RSiteSearch`.

Other `sos` functions provide summaries with one line for each package, support the union and intersection of `findFn` objects, and write the results to an Excel file with three sheets: (1) `PackageSum2`, which provides an enhanced summary of the packages with matches, (2) the `findFn` table itself, and (3) the call used to produce it.

Three examples are considered below: First we find a dataset containing a variable `Petal.Length`. Second, we study R capabilities for splines, including looking for a function named `spline`. Third, we search for contributed R packages with capabilities for solving differential equations.

Finding a Variable in a Data Set: `Petal.Length`

(Chambers, 2009, pp. 282-283) uses a variable `Petal.Length` from a famous Fisher data set but without naming the dataset nor indicating where it can be found nor even if it exists in any contributed R package. The sample code he provides does not work by itself. To get his code to work to produce his Figure 7.2, we must first obtain a copy of this famous data set in a format compatible with his code.

To look for this data set, one might first try the `help.search` function. Unfortunately, this function returns nothing in this case:

```
> help.search('Petal.Length')
No help files found ...
```

When this failed, many users might then try `RSiteSearch('Petal.Length')`. This produced 80 matches when it was tried one day (and 62 matches a few months later). `RSiteSearch('Petal.Length', 'function')` will identify only the help pages. We can get something similar and for many purpose more useful as follows:

```
> library(sos)
> PL <- findFn('Petal.Length')
```

`PL` is a `data.frame` of class `findFn` identifying all the help pages in Jonathan Baron's data base matching the search term. An alias for "findFn" is "???", and this same search can be performed as follows:

```
> PL <- ???Petal.Length
```

This `data.frame` has columns `Count`, `MaxScore`, `TotalScore`, `Package`, `Function`, `Date`, `Score`, `Description`, and `Link`. `Function` is the name of the help page, not the name of the function, as multiple functions may be documented on a single help page, and some help pages document other things such as data sets. `Score` is the index of the strength of the match. It is used by Baron's search engine to decide which items to display first. `Package` is the name of the package containing `Function`. `Count` gives the total number of matches in `Package` found in this `findFn` call. By default, the `findFn` object is sorted by `Count`, `MaxScore`, `TotalScore`, and `Package` (to place the most important `Package` first), then by `Score` and `Function`.

The summary method for such an object prints a table giving for each `Package` the `Count` (number of matches), `MaxScore` (max of `Score`), `TotalScore` (sum of `Score`), and `Date`, sorted like a Pareto chart to place the `Package` with the most help pages first:

```
> summary(PL)
```

```
Total number of matches: 27
Downloaded 27 links in 14 packages.
Packages with at least 1 match using search
                                pattern 'Petal.Length':
Package Count MaxScore TotalScore      Date
yaImpute    8         1           8 2009-08-16
<...>
datasets    1         2           2 2009-07-09
<...>
```

One of the listed packages is `datasets`. Since it is part of the default R distribution, we decide to look there first. We can select that row of `PL` just like we would select a row from any other `data.frame`:

```
> PL[PL$Package=='datasets', 'Function']
[1] iris
```

Problem solved in less than a minute! Any other method known to the present authors would have taken substantially more time.

Finding Packages with Spline Capabilities

Almost four years ago, the lead author of this article decided he needed to learn more about splines. A literature search began as follows:

```
RSiteSearch('spline')
```

(using the `RSiteSearch` function in the `utils` package). While preparing this manuscript, this command identified 1526 documents one day. That is too many. It can be restricted to functions as follows:

```
RSiteSearch('spline', 'fun')
```

This identified only 739 one day (631 a few months earlier). That's an improvement over 1526 but is still too many. To get a quick overview of these 739, we can proceed as follows:

```
splinePacs <- findFn('spline')
```

This downloaded a summary of the 400 highest-scoring help pages in the 'RSiteSearch' data base in roughly 5-15 seconds, depending on the speed of the Internet connection. To get all 739 matches, increase the `maxPages` argument from its default 20:

```
splineAll <- findFn('spline', maxPages=999)
```

The `print` method for a `findFn` object displays the result as a table in a web browser.

If we want to find a function named `spline`, we can proceed as follows:

```
selSpl <- (splineAll[, 'Function'] == 'spline')
splineAll[selSpl, ]
```

This has 0 rows, because there is no help page named `spline`. This does not mean that no function with that exact name exists, only that no help page has that name. To find a function with that exact name, try `findFn('spline')`. This produced one match for a function named 'regspline'.

To look for functions whose name includes the characters 'spline', we can use `grepFn`:

```
grepFn('spline', splineAll, ignore.case=TRUE)
```

This returned a `findFn` object identifying 78 help pages. The `print` method for an object of class `findFn` presents the result in a web browser. In this case, the sixth row is 'lspline' in the 'assist' package, which has a Score of 1. It is the fifth row in this table, because it is in the `assist` package, which had a total

of 34 help pages matching the search term, and this was the only one whose name matched the `grepFn` pattern.

To try to evaluate further the `splineAll` `findFn` object, we must first acknowledge that a table with 739 rows is too large to digest easily.

`summary(splineAll)` would tell us that the 739 help pages came from 191 different packages and display the first `minPackages` = 12 such packages. (If other packages had the same number of matches as the twelfth package, they would also appear in this summary.)

A more complete view can be obtained in MS Excel format using the `writeFindFn2xls` function:

```
writeFindFn2xls(splineAll)
```

If either the `WriteXLS` package and compatible Perl code are properly installed or if you are running Windows with the `RODBC` package, this produces a '*.xls' Excel file with three sheets:

The `PackageSum2` sheet contains information on locally installed packages not available from `summary`.

The `findFn` sheet contains the search results.

The `call` sheet gives the call to `findFn` that generated these search results.

If `WriteXLS` cannot produce an Excel file with your installation, it will write three *.csv files. (NOTE: Users who do not have MS Excel may like to know that Open Office Calc can open a standard '*.xls' file and can similarly create such files (Openoffice.org, 2009).)

The `PackageSum2` sheet (or file) is created by the `PackageSum2` function, which adds information from installed packages not obtained by `findFn`. This includes the package title and date, plus the names of author and maintainer, the date packaged, the number of help pages in the package, and the name(s) of any vignettes. This can be quite valuable in prioritizing packages for further study. I'd rather learn how to use a package being actively maintained than one that has not changed in five years. Similarly, I might prefer to study a capability in a larger package than a smaller one, because the rest of the package might provide other useful tools or a broader context for understanding the capability of interest.

For packages not already installed, the standard `install.packages` function in the `utils` package can be used. To make it easier to do this, the `sos` package includes a `installPackage` function, which checks all the packages in a `findFn` for which the number of matches exceeds a second argument `minCount` and installs any of those not already available locally; the default `minCount` is the square root of the largest Count. Therefore, the results from `PackageSum2` and the `PackageSum2` sheet of `writeFindFn2xls` will typically contain more information after running `installPackages` than before.

To summarize, two lines of code gave us a very powerful summary of spline capabilities in contributed R packages:

```
splineAll <- findFn('spline', maxPages=999);
writeFindFn2xls(splineAll)
```

The resulting splineAll.xls file can help establish priorities for further study of the different packages and functions. An analysis of this nature almost four years ago led the lead author to the *fda* package and its companion books, which further led to a collaboration that has produced joint presentations at three different conferences and a joint book (Ramsay et al., 2009).

Combining Search Results to Find Functions to Solve Differential Equations

The lead author of this article recently gave an invited presentation on "Fitting Nonlinear Differential Equations to Data in R" (Graves et al., 2009). A key part of preparing for that presentation was a search of contributed R code, which proceeded roughly as follows:

```
de <- findFn('differential equation')
des <- findFn('differential equations')
de. <- de | des
```

The object *de* has 53 rows, while *des* has 105. If this search engine were simply searching for character strings, the first would be larger than the second rather than the other way around. The last object *de.* is the union of the other two; *|* is an alias for *unionFindFn*. The *de.* object has 124 rows, which suggests that the corresponding intersection must have $(53+105-124) = 34$. This can be confirmed via `nrow(de & des)`.

To make everything in *de.* locally available, we can use `installPackages(de., minCount=1)`. This installed all referenced packages except *rmutil* and a dependency *Biobase*, which were not available on CRAN but are included in Jonathan Baron's "R site search" data base.

Next, `writeFindFn2xls(de.)` produced a file *de..xls* in the working directory (identifiable via `getwd()`).

The *PackageSum2* page of that Excel file provided a quick summary of packages with matches, sorted to put the package with the most matches first. In this case, this first package was *deSolve*, which provides, "General solvers for initial value problems of ordinary differential equations (ODE), partial differential equations (PDE) and differential algebraic equations (DAE)". This is clearly quite relevant to the subject. The second package was *PKfit*, which is "A

Data Analysis Tool for Pharmacokinetics". This may be too specialized for general use. I therefore would not want to study this first unless my primary interest here was in pharmacokinetic models.

By studying this summary page in this way, I was able to decide relatively quickly which packages I should consider first. In making this decision, I gave more weight to packages with one or more vignettes and less weight on those where the 'Packaged' date was old, indicating that the code was not being actively maintained and updated. I also checked the conference information to make sure I didn't embarrass myself by overlooking a package authored or maintained by another invited speaker.

Discussion

In sum, we have found *findFn* in the *sos* package to be very quick, efficient, and effective for finding things in contributed packages. The *grepFn* function helps quickly look for functions (or help pages) with particular names. The *unionFindFn* and *intersectFindFn* (especially via their *|* and *&* aliases) can be quite useful where a single search term seems inadequate; they make it easy to combine multiple searches to produce something closer to what is desired. An example of this was provided with searching for both "differential equation" and "differential equations".

Finally, the "PackageSum2" sheet of an excel file produced by `writeFindFn2xls` is quite valuable for understanding the general capabilities available for a particular topic. This could be of great value for authors to find what is already available so they don't duplicate something that already exists and so their new contributions appropriately consider the contents of other packages.

The *findFn* capability can also reduce the risk of "the researcher's nightmare" of being told after substantial work that someone else has already done it.

Acknowledgments

The capabilities described here extend the power of the *RSiteSearch* search engine maintained by Jonathan Baron. Without Prof. Baron's support, it would not have been feasible to develop the features described here. Duncan Murdoch, Marc Schwarz, Dirk Eddelbuettel and Gabor Grothendiek and anonymous referees contributed suggestions for improvement, but of course can not be blamed for any deficiencies. The collaboration required to produce the current *sos* package was greatly facilitated by R-Forge (R-Forge Team, 2009). The *sos* package is part of the *RSiteSearch* project hosted there. This project also includes code for a Firefox extension to simplify the process of finding information about R from within Firefox.

Spencer Graves
 President and Chief Operating Officer
 Structure Inspection and Monitoring
 San Jose, CA
 email: spencer.graves@prodsyse.com

Sundar Dorai-Raj
 Google
 Mountain View, CA
 email: sdorairaj@google.com

Romain François

email: romain.francois@dbmail.com

Bibliography

J. Baron. R site search.
<http://finzi.psych.upenn.edu/search.html>,

September 2009.

J. Chambers. *Software for Data Analysis: Programming with R*. Springer, New York, 2009.

S. Graves, G. Hooker, and J. Ramsay. Fitting nonlinear differential equations to data in r. http://stat.sfu.ca/~dac5/workshop09/Spencer_Graves.htm. August 2009. conference presentation.

Openoffice.org. *Open Office Calc*. Sun Microsystems, California, USA, 2009. URL <http://www.openoffice.org>.

R-Forge Team. R-forge. <http://r-forge.r-project.org>, September 2009.

J. Ramsay, G. Hooker, and S. Graves. *Functional Data Analysis with R and MATLAB*. Springer, New York, 2009.